



Cost landscapes by paradigm: Classical ML · QML (non-variational) · QML (variational)

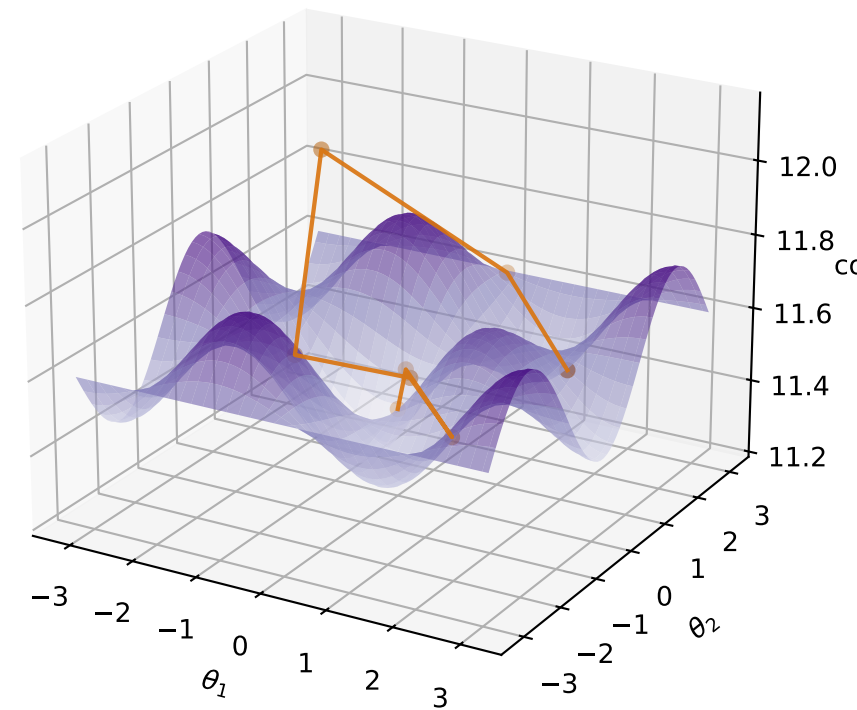
(a) Classical ML — Logistic Regression
convex basin · accuracy = 1.00



(b) QML (non-variational) — QSVM
flat · no trainable params · accuracy = 0.33



(c) QML (variational) — VQC
rugged / near-flat · accuracy = 0.50



(a) real logistic loss on Iris. (b) flat by design: non-variational QML has no parameter-training landscape; height = 1 - accuracy. (c) surface RECONSTRUCTED from collected convergence stats; trajectory z = real measured cost, (theta1, theta2) illustrative. No QPU was used.